

Suraj Kumar Rawani

+917061205601 ◇ rawanisuraj1919@gmail.com ◇ New Ashok Nagar, New Delhi -110096

◇ [LinkedIn](https://www.linkedin.com/in/suraj-kumar-rawani-0483b7298/) : <https://www.linkedin.com/in/suraj-kumar-rawani-0483b7298/>

◇ [Github](https://github.com/SurajRawani1919) : <https://github.com/SurajRawani1919>

SUMMARY

AI/ML Engineer with 1+ year of hands-on experience building **Generative AI applications, RAG pipelines, LLM evaluation workflows, backend services, and data processing systems**. Strong in **Python, NLP, LLMs, prompt engineering, embeddings, LangChain, Hugging Face, FAISS, SQL, Pandas, and NumPy**. Experienced in developing document-based question-answering systems, vector-based retrieval pipelines, interactive AI applications, and automated LLM evaluation workflows. Hands-on experience with **Git/GitHub, Docker, FastAPI, Streamlit, and Databricks**.

SKILLS

- **Programming & Data:** Python, SQL, Pandas, NumPy
- **Generative AI & NLP:** Generative AI, LLMs, NLP, RAG, Prompt Engineering, Embeddings, LLM Evaluation, Agentic AI
- **AI Frameworks & Libraries:** LangChain, Hugging Face, FAISS, Scikit-learn, Streamlit
- **Backend & APIs:** FastAPI, REST APIs
- **Databases & Retrieval:** SQL, FAISS, Vector-Based Retrieval, Vector Databases
- **Development & Deployment:** Git, GitHub, Docker, CI/CD, Linux, Shell/Bash, YAML
- **Data & ML:** Machine Learning, Deep Learning, Data Pipelines, Data Processing, Computer Vision
- **Tools:** Databricks

PROJECTS

NVIDIA SysBench – LLM Instruction Following Benchmark

- Designed and validated multi-turn, constraint-driven evaluation tasks to benchmark large language models on instruction-following, reasoning stability, and rule prioritization.
- Engineered deterministic system prompts with atomic, verifiable constraints and ensured exact alignment with turn_metadata for automated and human evaluation.
- Created aligned and adversarial (misaligned) dialogue scenarios to test safe refusal behavior, long-context retention, and multi-turn stability.
- Evaluated model performance using CSR, ISR, and SSR metrics, improving task robustness against instruction drift and hallucinations.
- Utilized NVIDIA CIF (Streamlit validation tool) and Turing ChatHub to validate notebooks, debug constraint violations, and iterate on LLM behavior.

Key Skills & Keywords:

LLM Evaluation & Benchmarking, Instruction Following (IF), Prompt Engineering (System Prompts), Multi-Turn Reasoning, Constraint Satisfaction (CSR/ISR/SSR), Safe Refusal & Alignment, LLM Validation Pipelines, NVIDIA & Turing Benchmarks

RAG-based LLM Application

- Built a Retrieval-Augmented Generation (RAG) system combining Large Language Models (LLMs) with a vector database for efficient contextbased responses.
- Implemented document ingestion, embedding generation, and retrieval pipeline to enhance model accuracy with external knowledge.
- Integrated with Hugging Face models and APIs for natural language understanding.
- Deployed locally with Streamlit to create an interactive chat interface.
- Utilized Python, LangChain, Hugging Face, FAISS, and Streamlit for development

Handshake AI (Contractual) — Project Dynamo | AI Evaluation

- Worked on **Terminal-Bench 2 (TB2)** tasks under the Handshake AI Project Dynamo initiative.
- Designed and developed **terminal-based coding and software engineering tasks** for AI model evaluation.
- Created task environments using **Docker**, including Dockerfile, task.toml, instructions, and test configurations.
- Used **Harbor CLI** to build, validate, and submit benchmark tasks.
- Developed **automated test scripts and evaluation logic** to verify whether AI agents successfully completed assigned tasks.
- Designed tasks focused on **coding, debugging, file manipulation, CLI workflows, and instruction following**.
- Reviewed and reworked tasks based on **evaluation feedback and benchmark guidelines**.
- Performed quality checks to ensure tasks were **clear, reproducible, diverse, and appropriately challenging**.
- Contributed to **AI agent benchmarking and evaluation** by testing task execution in containerized environments.

INTERNSHIP/EXPERIENCE

AI/ML Engineer

Jupiter AI Labs | Noida 63 | Sep 2025 - Present

- Develop and train machine learning models for real-world AI applications using Python and modern ML frameworks.
- Design and maintain data pipelines for **data ingestion, preprocessing, transformation, feature engineering, and model evaluation** using Python, SQL, and Pandas.
- Implement structured data-processing workflows for preparing large datasets for machine learning training and evaluation.
- Build backend services and automation workflows for AI applications using modern development tools and APIs.
- Collaborate with engineering teams to test model components, validate outputs, and document technical workflows.
- Utilize **Docker, Git/GitHub, and Databricks** to maintain reproducible development environments and version-controlled ML workflows.
- Apply data analysis and visualization techniques to validate model performance and support data-driven development.
- Contribute to improving AI development workflows through data-processing and model-evaluation optimization.

Data Scientist Intern

Regex Software Services | Jaipur | Nov 2024 – May 2025

- Completed training in Data Science with hands-on experience in data collection, preprocessing, and analysis.
- Completed hands-on training in **Data Science, data collection, preprocessing, analysis, and machine learning workflows**.
- Worked on data pipelines to clean, transform, and prepare datasets for machine learning applications.
- Used **Python, SQL, and analytics tools** to manage and analyze structured datasets.
- Applied data analysis and visualization techniques to extract insights from datasets.

EDUCATION

Master of Computer Applications, CHANDIGARH UNIVERSITY (PUNJAB)

Jun'25 — Expected 2027

Bachelor of Computer Application, SURESH GYAN VIHAR UNIVERSITY (JAIPUR)

Jun'22 — Apr '25

CERTIFICATIONS

Databricks Fundamentals Accreditation, Databricks

Advanced LLMs with Retrieval Augmented Generation (RAG)